


A U.S. lawyer used ChatGPT to help draft court filings and ended up **citing** entirely fake legal cases. When the opposing side challenged the citations, the lawyer claimed not to realize that ChatGPT was a generative language tool rather than a reliable legal database.


A federal judge responded with a standing order obliging anyone appearing before the court to attest that AI was not used to draft court filings or flag any language drafted by AI to be checked for accuracy.



Smart search recommends putting glue on pizza

Google's "AI Overview" feature **went viral** because of the bizarre responses it generated. For example, the AI suggested adding non-toxic glue to pizza sauce to make cheese stick (some users actually **tried** doing that!). Presumably, these bogus answers were derived from misinterpreted satirical or forum content that LLM used to generate the answer.

The absurd responses sparked questions about the reliability of Google's generative AI approach, leading the company to announce revisions to its system.



Source: Peter Yang's [account](#) on X

How to reduce AI hallucinations

AI hallucinations are a tricky beast. LLMs make things up so confidently that detecting fabricated information can be difficult. The good news is that there are several things you can do to reduce AI hallucinations.

Ground the outputs in trusted data. You can use [Retrieval-Augmented Generation \(RAG\)](#) to ground LLM outputs in trusted data sources, such as company policies or documents. For example, in a RAG customer support setup, the AI first searches the company's support documentation before replying to a customer, ensuring its response is accurate and consistent with the latest company guidelines.

Use curated datasets. For high-risk use cases like healthcare, you can fine-tune your model on curated or domain-specific corpora.

Adjust a prompt. You can set clear constraints to ground the response only on the provided context and to admit uncertainty – e.g., “Answer only using the provided context. If the answer is not in the context, say 'Not found.'” You can also encourage the model to admit uncertainty – e.g., “If unsure, say you don't know.”

Add verification and validation layers. You can use [LLM-as-a-judge](#) to validate and fact-check the outputs. For regulated or high-stakes domains, add human validation workflows.

Evaluate and monitor. Finally, you need [LLM evaluations](#) to ensure the outputs of your AI app are accurate, safe, and aligned with user needs. Once live, continuously monitor production outputs to detect new failure modes.