



Jukka Niiranen's Post



Jukka Niiranen

6mo



It took me 3 prompts with GPT-5 to get a hallucination from it. How about you?

I was reading the reviews from AI influencers who had preview access to the new GPT-5 models. Namely 2 people whose writings I highly value. I paid attention to what both said about hallucinations:

Simon Willison: "In my own usage I've not spotted a single hallucination yet."

Ethan Mollick: "...these systems still make errors and generate hallucinations that humans need to check (although I did not spot any major issues in my own use)."

Okay, so opening ChatGPT didn't yet show GPT-5 to me. But luckily Microsoft was really quick this time (well done! 🙌) and the "Try GPT-5" button showed up there on first load of the web UI.

Given how the classic "how many R's in strawberry" is surely in the training data a million times for these newest models, I had to try something different. And what do you know: instant failure. 🌟

Yeah, the total number of R's was right. However, in particular the claim about an R in "fields" showed how LLMs are still LLMs. Missing the other R in "forever" made it a success score of 33% in identifying characters from a string of 25 characters.

Evidently GPT-5 should be smart in choosing what model to use for which problem. Maybe it was overconfident about strawberries and decided to save up on GPU resources, picking a mini model.

It doesn't matter for the users. The bottom line remains the same: the failure modes of AI assistants today are something which normal folks will not be able to comprehend. They will fall for the sweet talking AI that sounds very smart, yet doesn't have intelligence. And thus, when pressured to use more AI in their work tasks to align with top leadership goals, errors will spread even further.

Yeah, I will keep using GPTs daily. When you know how the models can fail and stay on your toes about the responses potentially being BS, it's an improvement over not having LLMs.

But it is what it is. Still GPT...

99 · 62 Comments

Like

Comment

Share

George Hepworth

6mo

Part 2.

And that leads to the second point:

It's a process, not a quiz. You don't just accept the 1st response at face value. You have to be prepared to react to the response appropriately as well. You are, as a friend of mine likes to point out, the adult in the room, not ChatGPT.

I'm quite sure ChatGPT would be obsequious if you pointed out the error in the reported r count. And it would get it right on the second or third attempt.

Of course, we already know the answer to the question before asking, so it's easy to spot and address the error.

The problem is that many times, the context and content of a response is so complex that we may not recognize flaws until it's too late.

And that leads to a conclusion about validation.

If you fail to validate any response from an AI, the consequences are on you, not the AI.

Refer back to the observation about being the adult in the room.

Like · Reply | 1 Reaction

George Hepworth

6mo

Two observations about using ChatGPT as an assistant.

First, Prompt Engineering has a lot to do with the quality of the responses generated. As the old sayings goes, "Garbage in, Garbage Out."

Before you leap to the claim that ChatGPT obviously has been given garbage training, no. That's not what I mean.

I mean the quality of the prompt you use is closely related to the quality of the output you get back. It's like asking a friend who happens to be an accountant to "explain the tax laws". That prompt is likely to elicit a good deal of fluff and nonsense before you get to useful information.

As someone else in the thread pointed out, the original prompt carried at least one level of assumption. You are interested in both capital and lower case R. It's safe to assume that a human will share your interpretation of what counts. Well, at least those of us educated prior to the age of texting with our thumbs. 🙄 If you want an explicit count of both, a good prompt would be specific on that point. ChatGPT was trained by thumb-